


# Analysis of computer clusters

Christian Faccio

March 28, 2025

# Cluster Characteristics



- **One master node**  
connected to the Internet through NAT connection
- **Two worker nodes**  
connected to the master node with Internal Network (not directly to Internet)
- 2GB of RAM, 2 CPUs, 25GB of storage each

- ❶ Configuration of the master node
  - Adapters
  - Port Forwarding
  - (SSH connection to the host)
  - Network configuration
- ❷ Creation of the worker nodes
- ❸ DHCP and DNS configuration
- ❹ Configuration of the Master node as Gateway
- ❺ SSH on worker nodes
- ❻ Creation of a distributed file system

# Containers' steps

- 1 Creation of a Dockerfile
- 2 Creation of the docker-compose.yaml file
- 3 Start the containers

```
docker-compose up -d
```

```
ssh -i ssh_keys/id_rsa -p <port> user@localhost
```

# Dockerfile

```
FROM ubuntu:latest

# Set non-interactive mode for apt-get
ENV DEBIAN_FRONTEND=noninteractive

# Install required packages
RUN apt-get update && apt-get install -y \
    openssh-server ssmc iputils-ping \
    sysbench stress-ng iostat3 iperf3 \
    netcat-openbsd wget unzip htop \
    epich vim \
    openmpi-bin openmpi-common openmpi-doc libopenmpi-dev \
    sudo \
    && rm -rf /var/lib/apt/lists/*

# Create SSH folder and set correct permissions
RUN mkdir -p /var/run/ssh /home/user/.ssh /shared \
    && chmod 700 /home/user/.ssh

# Create a new user 'user' with a home directory
RUN useradd -m -s /bin/bash user \
    && echo "user:userpassword" | chpasswd \
    && echo "user ALL=(ALL) NOPASSWD:ALL" >> /etc/sudoers

# Ensure SSH is configured for user
RUN sed -i 's/#PermitRootLogin prohibit-password/PermitRootLogin no/' /etc/ssh/sshd_config \
    && sed -i 's/#UsePrivilegeSeparation no/UsePrivilegeSeparation no/' /etc/ssh/sshd_config \
    && sed -i 's/#PubkeyAuthentication yes/!PubkeyAuthentication yes/' /etc/ssh/sshd_config \
    && sed -i 's/#AuthorizedKeysFile .ssh/authorized_keys .ssh/authorized_keys/' /etc/ssh/sshd_config

# Copy SSH keys for user (passwordless login)
COPY ssh_keys/id_rsa.pub /home/user/.ssh/authorized_keys
COPY ssh_keys/id_rsa /home/user/.ssh/id_rsa

# Set correct permissions for SSH keys (user)
RUN chmod 600 /home/user/.ssh/id_rsa /home/user/.ssh/authorized_keys \
    && chown -R user:user /home/user/.ssh

# Expose SSH port
EXPOSE 22

# Switch to user
USER user
WORKDIR /home/user

# Start SSH service correctly with host key generation
CMD sudo ssh-keygen -A && sudo /usr/sbin/sshd -D -o && sudo chown -R user:user /shared &
```

Figure: Dockerfile

```
services:
  master:
    build: . # Use the Dockerfile in the current directory to build the image
    container_name: master # Set the container name to 'master'
    networks:
      - my_network # Attach this container to the custom bridge network
    deploy:
      resources:
        limits:
          cpus: '2' # Restrict the container to use a maximum of 2 CPU cores
          memory: 2G # Limit the container's memory usage to 2GB
        ports:
          - "2220:22" # Map port 2220 on the host to port 22 inside the container (SSH access)
      volumes:
        - shared_volume:/shared # Mount shared volume for data exchange between container
        - ./ssh_keys:/root/.ssh # Mount pre-generated SSH keys for passwordless access
      tmpfs:
        - /shared:nodev=777 # Create a temporary filesystem at /shared with full permissions
  node1:
    build: . # Use the same Dockerfile for worker node
    container_name: node1 # Set the container name to 'node1'
    networks:
      - my_network # Attach to the same network as the master node
    deploy:
      resources:
        limits:
          cpus: '2'
          memory: 2G
        ports:
          - "2221:22" # Assign a different SSH port for node1
      volumes:
        - shared_volume:/shared # Mount shared volume to enable data sharing
        - ./ssh_keys:/root/.ssh # Use the same SSH keys for passwordless login
      tmpfs:
        - /shared:nodev=777 # Temporary shared filesystem with full permissions
  node2:
    build: . # Build using the same configuration
    container_name: node2 # Name this container 'node2'
    networks:
      - my_network
    deploy:
      resources:
        limits:
          cpus: '2'
          memory: 2G
        ports:
          - "2222:22" # Assign another unique SSH port for node2
      volumes:
        - shared_volume:/shared
        - ./ssh_keys:/root/.ssh
      tmpfs:
        - /shared:nodev=777

networks:
  my_network:
    driver: bridge # Use a bridge network for inter-container communication

volumes:
  shared_volume:
    driver: local # Use a local volume for persistent shared storage
```

Figure: docker-compose.yml

Different tests have been performed to measure the performances of the clusters:

- **HPCC**: Tests computation, memory access, and communication efficiency.
- **Iperf3**: Tests network performance between nodes.
- **Stress-ng**: Tests CPU, memory, and I/O performance under stress.
- **Sysbench**: Tests CPU and memory performance under stress.
- **IOzone**: Tests disk I/O performance.

```
mpirun -np 4 -hostfile hosts hpcc
```

Test	Unit	VMs	Containers
MPIRandomAccess	GUP/s	0.003	0.01
PTRANS (Wall)	s	0.550	0.42
StarDGEMM	Gflop/s	2.909	3.95
StarSTREAM Copy	GB/s	23.20	21.66
MPIFFT	Gflop/s	2.133	8.43
Avg. Ping Pong Bandwidth	GB/s	4.756	17.148
HPL	Gflop/s	10.64	12.32

Table: Main results of the HPCC tests

Overall, containers provide a more efficient execution environment for most HPCC workloads, especially in **memory access**, **computational efficiency**, and **communication latency**.

```
iperf3 -s #on the server
```

```
iperf3 -c <server_ip> #on the client
```

Category	Nodes	Transfer(GB)	Bitrate(GB/s)
VMs	Master-Node1	2.09	1.79
	Node1-Node2	2.15	1.84
Containers	Master-Node1	142	122
	Node1-Node2	140	120

**Table:** Main results of the Iperf3 tests categorized by type

Overall, containers are more efficient, with significantly **higher transfer rates** and **fewer network issues** compared to VMs.



# Stress-ng

```
mpirun -hostfile hosts -np 4 stress-ng -cpu 2 -timeout 60s
```

```
mpirun -hostfile hosts -np 4 stress-ng -vm 2 -vm-bytes 1G
```

```
mpirun -hostfile hosts -np 4 stress-ng -io 2 -timeout 60s
```

Test	Bogo ops		Bogo ops/s	
	VM	Containers	VM	Containers
CPU	21532	24523	358.71	408.43
Memory	760052	7099658	12634.09	118275.80
I/O	2799634	505508	46659.86	8425.03

Table: Main results of the Stress-ng tests

Overall, containers demonstrate **better CPU and memory handling**, while VMs outperform containers in **disk I/O operations**.


```
mpirun --hostfile hosts -np 4 sysbench cpu  
--cpu-max-prime=20000
```

Test	Ops/s		Avg. Latency	
	VM	Containers	VM	Containers
CPU	2,198.45	3,303.00	0.46	0.31
Memory	19,327.09	41,471.36	0.05	0.03


**Table:** Main results of the Stress-ng tests (max values)

Overall, containers demonstrate better CPU and memory handling, with **higher throughput** and **lower latency** compared to VMs.


## VM Read




## Containers Read



## VM Write



## Containers Write



`iozone -+m /shared/machines.txt -f /shared/testfile -a -R -O`

Containers generally provide better **performance** and **efficiency** for high-performance computing, networking, and I/O workloads. VMs still offer **stronger isolation**, but for resource-intensive tasks with heavy inter-process communication, containers tend to be **faster** and more **lightweight**.